

NN+NLP Project

Team members

- Jakub Kuciński
- Wojciech Fica
- Jakub Dworzański

Project title

Musical notes generation

Description

We build NN models that generate subsequent musical notes when given some (short) initial sequence of notes.

What we wanted to achieve and why:

- always wanted to be a musician but didn't have the talent
- see if we can produce good quality music
- it was a new and challenging topic
- AI should be able to generate music

Was there any challenge?

- Seemed similar to Language Modelling
- Yet different: music structure != language structure
 - language sentence == a tree of words
 - music == a flow of notes + ...

Data

[Lakh Pianoroll Dataset](#) is a collection of 174,154 multitrack pianorolls derived from the Lakh MIDI Dataset (LMD).

Data preparation:

How to turn



into tokens?

- Brainstormed a few approaches used in the field.
- Chose the simplest one.


Sample encoding:

xxsep d34 n44 d2 n39 d6 n32 d2

xxsep d1 n44 d2 n33 d3

xxsep d2 n46 ...

that corresponds to



Models

RNN-based

- LSTM:
 - written and trained from scratch
 - embedding_dim = 128, lstm_size = 256, num_layers=2, ...

Transformer-based:

- GPT2 (LMHeadModel)
 - wanted to finetune a model from a library
 - the tricky part: model was not trained on musical notes ...
 - so decided to train a mini GPT2 from scratch instead
 - n_ctx = 128, n_embd = 64, n_layer = 3, ...

Success evaluation

Empirical - if model generates reasonably fine-sounding and fairly diverse songs based on different initial notes we would call it a success.

How to evaluate a notes-generating model?

1. Does the generated music sound good? - YES
2. Is the model able to generate music or just remembered a few pieces? - YES
3. Perplexity:

- LSTM perplexity:
- GPT2 perplexity:

Conclusions

What have you learned?

- data preparation may get tricky and might require domain specific knowledge;
- models take a long time and a lot of resources to train.

What was good or bad?

- good: the results were cool;
- bad: didn't have enough time/resources to finish training the models.

What could have been different?

- encode more information when encoding notes to tokens.

What could you do next?

- expand it to handle multiple instruments.

Examples

Examples with generated music in data/ directory.

Literature

[Deep Learning Music Generation](#)

[MuseNet](#)

[Compound Word Transformer: Learning to Compose Full-Song Music cover Dynamic Directed Hypergraphs](#)